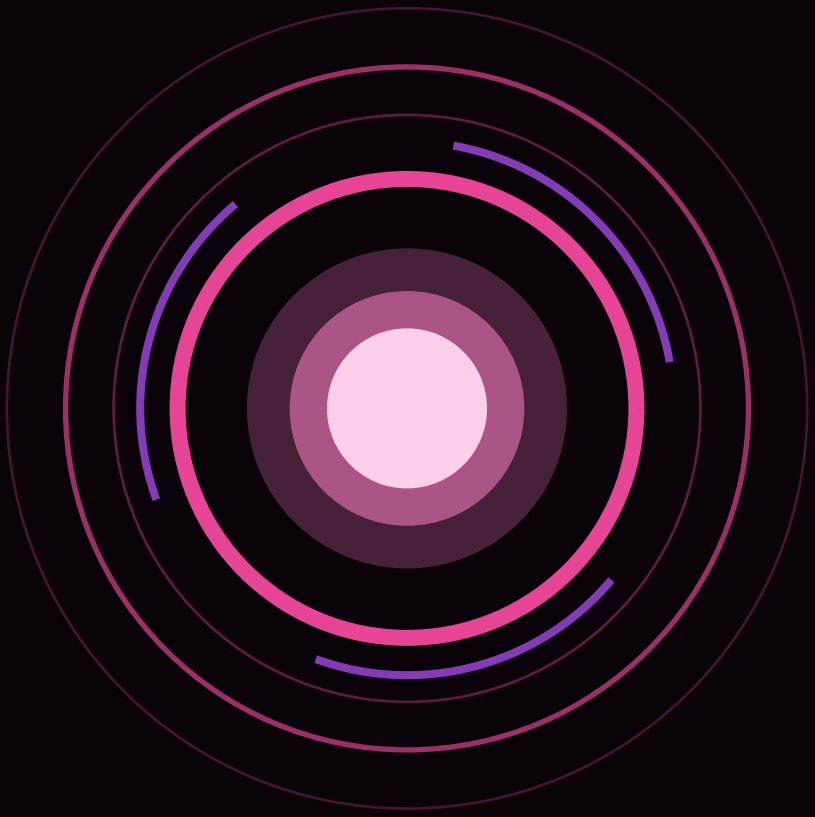

K E E N A

Your own AI companion, running on your PC

Setup guide: requirements, how she works, step-by-step install

Needs a GPU with 8 GB VRAM (NVIDIA recommended)

Contents

1	**What is Keena?**	Overview and everything she can do
2	**System requirements**	Minimum and recommended hardware; what isn't supported
3	**How Keena works**	From your voice to her reply, and where each step runs
4	**Step-by-step setup**	Twelve steps, each with a verification check
5	**Things to try**	Example phrases
6	**Updating and uninstalling**	Keeping Keena current, or removing her cleanly
7	**Safety, troubleshooting and FAQ**	Built-in safeguards and fixes for common problems
8	**Command reference**	Every command in one table

Online version: the same guide, always up to date, is in the repository at **github.com/WeDigiStudio/Keena-Personal-Assistant** (docs/SETUP.md).

What is Keena?

Keena is a personal AI companion that lives on your own Windows PC. You talk to her with your voice, in English, Tamil or Tanglish, and she answers in English, in a warm, natural voice. She chats like a friend, remembers what you ask her to remember, and actually controls your PC: she opens apps, plays songs on YouTube, searches Amazon, writes and sends emails in Gmail, answers questions from the web, and you can even control her from your phone.

Her brain is **Qwen 3 8B**, an open AI model that runs **on your own graphics card** through Ollama. No ChatGPT subscription, no monthly fee. That is also why she needs a decent PC: read the requirements first.

What she can do

Area	Say something like...	What happens
Talk	"Keena, I'm making Maggi, want some?" / "I had a rough day."	A personal, playful reply; on a rough day she fires you up.
Apps and PC	"Open Chrome", "volume 30", "take a screenshot", "lock the PC"	Done instantly, no AI needed for common commands.
Music	"Play Believer", "pause", "next song", "close the YouTube video"	Plays the first video on YouTube.
Web answers	"What's the latest on ...?", "tell me more about ..."	Searches the web, reads the top pages, answers in 1-3 sentences with the source.
Recommendations	"Suggest a good Vijay movie"	Two or three real picks from the web, ending with "I think you'll love ...".
Shopping	"Search Amazon for a wireless mouse", "open the cheapest one", "what's the price?"	Drives her own Chrome window. Anything that buys or pays asks you first.
Email	"Draft a mail to my CTO asking about the project, and send it"	Gmail opens, the mail types itself, she asks "Send it?" and sends only on yes.
Memory	"Remember that Arjun is my CTO", "what do you know about me?", "forget that"	Saved on your PC, only when you ask.
Reminders	"Remind me to call mom at 6 PM", "set a timer for 10 minutes"	She announces them on time.
Phone	Open the Keena page on your phone, tap once, then "Hey Keena..."	Same Keena, controlling your PC from your phone over Wi-Fi.

System requirements

The one that matters most: a graphics card with at least 8 GB of VRAM, ideally NVIDIA. Keena's brain (Qwen 3 8B, about 5.6 GB) and her hearing (Whisper, about 0.5-1 GB) both live in the graphics card's memory at the same time. On the 8 GB card she was built on, about 6.8 GB of the 8 GB is in use. With less VRAM she either doesn't fit or becomes very slow.

Part	Minimum	Recommended
Operating system	Windows 10 or 11, 64-bit	Windows 11
Graphics card (GPU)	NVIDIA RTX with 8 GB VRAM: RTX 2070 / 2080, RTX 3060 (8 GB), 3060 Ti, 3070, RTX 4060, 4060 Ti. AMD Radeon 8 GB+: experimental (below)	NVIDIA with 12 GB or more: RTX 3060 12 GB, RTX 4070 / 4070 Super, or better
System memory (RAM)	16 GB	32 GB
Processor (CPU)	6 cores (Intel Core i5 10th gen / AMD Ryzen 5 3600 or newer)	8 cores or more
Free storage	20 GB on an SSD	40 GB on an SSD
Internet	Needed: for her voice, web answers, Tamil, and the first setup downloads	Stable broadband
Audio	Any microphone + speakers	Headset or a good USB mic (she hears you better, and never hears herself)
Browser	Google Chrome installed	Google Chrome
Phone (optional)	Any Android or iPhone on the same Wi-Fi	Android with Chrome

NVIDIA or AMD?

GPU	Status	What changes
NVIDIA RTX, 8 GB+	Recommended, tested	Nothing: follow the guide as written.
AMD Radeon, 8 GB+ (supported by Ollama: mainly RX 6800 and newer, RX 7000 series)	Experimental, untested	Her brain runs on the Radeon through Ollama. Her hearing runs on the CPU: about 1-1.5 s slower per command (measured 1.4 s for a 5.5 s sentence, vs about 0.2 s on an RTX 4060). Follow the AMD notes in Steps 1, 5 and 6.
Intel Arc, integrated graphics	Not supported	Her brain would run on the CPU: far too slow.

Not supported

- **Mac and Linux:** Keena controls Windows itself (windows, keys, volume, apps), so she is Windows-only.
- **6 GB cards:** possible with the smaller model `qwen3:4b-instruct`, but she gets clearly less smart. Laptops with an 8 GB RTX card work.

How Keena works

Every time you speak, this happens, usually in one to two seconds:

#	Step	Where it runs
1	**Listening for her name.** The microphone is cut into short pieces; each is checked for "Keena".	Your PC
2	**Hearing.** Whisper (an open speech model) turns your voice into text on your graphics card in about 0.1-0.3 s.	Your PC (GPU)
3	**Tamil or Tanglish?** If Whisper isn't confident (it only knows English), the same audio goes to Google Gemini, which turns your Tamil into English. Adds about 2-3 s. Your English speech never leaves your PC.	Gemini (free key)
4	**Instant commands.** Common commands (open, play, volume, reminders...) are matched by rules and done instantly, without any AI.	Your PC
5	**Thinking.** Everything else goes to Qwen 3 8B: a chat reply, or a decision to use a tool (web search, Gmail, Amazon, browser...).	Your PC (GPU)
6	**Safety checks in code.** Before anything is said or done, code checks it: no invented facts about you, no fake "I did it", no invented email addresses, and a yes from you before sending, buying or shutting down.	Your PC
7	**Speaking.** The reply is turned into Keena's voice sentence by sentence (Microsoft Edge's free voices), so she starts talking before the whole answer is written.	Microsoft (free, no key)
8	**The HUD.** A pink arc-reactor orb in the corner of your screen shows what she's doing and the sentence she's saying. Say "big screen mode" for full screen.	Your PC

What stays on your PC, and what goes online

Being honest about this matters, so here it is plainly:

Runs on your PC (private)	Uses the internet
Her brain (Qwen 3 8B), all reasoning and chat Speech recognition for English (Whisper) Wake word Your memories, profile and conversation history (in the data folder) All PC control, the HUD, the phone server	Her voice: the text she says is sent to Microsoft's free voice service Web answers and recommendations: a web search (LangSearch, free key) Tamil / Tanglish speech: the audio is sent to Google Gemini (free key) Whatever websites she opens for you (YouTube, Gmail, Amazon)

Step-by-step setup

Plan about 45-60 minutes, mostly downloads. Do the steps in order; each one ends with a check.

Step 1. Update your graphics driver

NVIDIA: install the latest driver from nvidia.com (or the NVIDIA app). Then open PowerShell and check your card and its memory:

```
nvidia-smi
```

You should see your RTX card and a memory total of about 8000 MiB or more.

AMD: install the latest AMD Software: Adrenalin Edition from amd.com/support. In Task Manager (Performance, GPU), your Radeon should show 8 GB or more, and it must be on Ollama's list of supported AMD GPUs.

Step 2. Install Python 3.11, Git and Google Chrome

- **Python 3.11** from python.org (version 3.11.x, not newer). In the installer, tick **"Add python.exe to PATH"**.
- **Git** from git-scm.com (default options).
- **Google Chrome** if you don't have it.

```
py -3.11 --version
```

Step 3. Install Ollama and download her brain

Install Ollama from ollama.com. It runs in the background (tray icon). Then download the model (about 5.2 GB):

```
ollama pull qwen3:8b
ollama run qwen3:8b "Say hi in five words"
ollama ps
```

`ollama ps` should show the model at **100% GPU**. If it says CPU, your card has too little free VRAM or the driver is old.

Step 4. Download Keena's code

Download the code from GitHub: **github.com/WeDigiStudio/Keena-Personal-Assistant**

```
cd $HOME
git clone https://github.com/WeDigiStudio/Keena-Personal-Assistant.git keena
cd keena
```

Step 5. Install her Python parts

Create a private Python environment inside the folder and install everything into it:

```
py -3.11 -m venv .venv
.venv\Scripts\activate
pip install -r requirements.txt
pip install nvidia-cublas-cu12 "nvidia-cudnn-cu12==9.*"
```

The last line gives Whisper its GPU libraries (**NVIDIA only: skip it on AMD**). If PowerShell refuses to run `activate`, run this once and try again:

```
Set-ExecutionPolicy -Scope CurrentUser RemoteSigned
```

From now on, always run `.venv\Scripts\activate` first in a new PowerShell window.

Step 6. Your settings file

```
copy .env.example .env
notepad .env
```

In Notepad, put your name after `JARVIS_USER=` and save. Everything else already has good defaults (voice: Ava; model: qwen3:8b; hearing on the GPU).

AMD: also change `WHISPER_DEVICE=cuda` to `WHISPER_DEVICE=cpu`. If you forget, she switches to the CPU by herself and says so at startup.

Step 7. Two free keys (recommended)

Key	What it unlocks	Where to get it
LANGSEARCH_API_KEY	Web answers, news, prices, movie and song recommendations	langsearch.com: sign up, create an API key
GEMINI_API_KEY	Understanding Tamil and Tanglish speech	aistudio.google.com: sign in, "Get API key"

Paste each key after its name in `.env`, with no spaces or quotes. Without them, Keena still works: she just can't search the web or understand Tamil. Free keys have daily limits that normal use doesn't reach. Never share your `.env` file or post your keys.

Optional: add your email domain after `KNOWN_EMAIL_DOMAINS=` (e.g. yourcompany.com) so addresses you say out loud are spelled right.

Step 8. Tell her about yourself

Create `about_me.txt` in the keena folder (Notepad is fine). One fact per line:

```
Name: Your Name
City: Chennai
Work: Founder of ...
Favourite food: Biryani
Best friend: ...
```

She reads this every time you talk about yourself, and never invents facts that aren't there. You can also say "Keena, remember that ..." any time; that's saved on your PC too.

Step 9. First start

Test in this order. The first start downloads Whisper (about 500 MB), so give it a minute.

```
python main.py --mic-test      # talk: the level should jump; Ctrl+C to stop
python main.py --text          # type to her: tests the brain and tools
python main.py                 # the real thing: say "Hey Keena"
```

You'll see a pink orb in the corner of the screen. Say **"Hey Keena"**, wait for the soft chime, then talk. After she answers, you can keep talking for a few seconds without saying her name again.

Step 10. Her own Chrome window: sign in once

The first time she opens a website, a **separate Chrome window** appears. It's her own Chrome profile (Chrome only allows a program to control a separate profile). Your normal Chrome, passwords and logins are untouched. Say **"Keena, open Gmail"** and sign in there once (and YouTube or Amazon if you want). The logins stay.

Step 11. Control her from your phone (optional)

Phone and PC on the same Wi-Fi. With Keena running, open a second PowerShell window in the folder:

```
.venv\Scripts\activate
python -m remote.pair
```

- Scan the QR code with your phone camera.
- Chrome says **"Your connection is not private"**: expected, because your PC made its own certificate. Tap **Advanced**, then **Proceed**. The connection is still encrypted.
- Allow the microphone, then Chrome menu: **Add to Home screen**. It opens like an app.
- Tap the reactor once (phones need one touch before they allow the mic and sound), then just say **"Hey Keena ..."**. It works while the page is open on screen.
- Windows may ask whether Python can use the network: allow it on private networks.

Step 12. Start her with Windows (optional)

```
powershell -ExecutionPolicy Bypass -File scripts\autostart.ps1
```

She then starts minimised when you log in, and restarts herself if she ever crashes. To turn it off, run the same line with `remove` at the end.

Things to try

Say	Keena
"Hey Keena, how are you?"	Chats back, in character.
"I'm so tired today."	"Hey ..., don't you remember who you are? Come on, you can do this..."
"Play Believer on YouTube" / "pause" / "volume 30"	Instant.
"Open Amazon and search for a gaming mouse, then open the cheapest one"	Searches, opens it, reads the price.
"Which one is the best?"	Compares real listings (rating, price, number of ratings).
"Suggest a good Rajinikanth movie"	Two or three real picks, then her favourite.
"Draft a mail to Arjun asking about the project and send it"	Gmail, typed live, "Send it?"
"Remind me to drink water in 30 minutes"	Reminds you on time.
"What do you know about me?" / "forget that"	Reads back or deletes your memories.
"Big screen mode" / "back to normal"	Full-screen HUD and back to the corner.
"Goodbye Keena"	Shuts her down.

Updating and uninstalling

Update to the latest version

```
cd $HOME\keena
.venv\Scripts\activate
git pull
pip install -r requirements.txt
```

Your `.env`, `about_me.txt` and data folder are not affected.

Uninstall

- Turn off autostart if you enabled it (Step 12, with `remove`).
- Delete the keena folder. This also deletes her memories and her browser profile.
- Optionally remove the model with `ollama rm qwen3:8b`, and uninstall Ollama.

Safety built in

- She **never sends an email, buys anything, shuts down or sleeps the PC without your yes**.
- She never types or presses keys into a terminal window (so she can't run commands by accident).
- She never invents email addresses, reminder times or facts about you: code checks every reply.
- She only remembers what you ask her to remember, and forgets on request.
- The phone page needs a secret key (in the QR code) and only works from your home network or your own Tailscale.
- Your conversations, memories and profile stay in the keena folder on your PC.

Troubleshooting

Problem	Fix
"Can't reach Ollama" when starting	Start Ollama (Start menu), wait a few seconds, run again.
She's very slow, or ollama ps shows CPU	Not enough free VRAM: close games and other GPU apps; check the card has 8 GB+.
She never hears "Hey Keena"	Run python main.py --mic-test; set MIC_DEVICE in .env; speak a little closer.
She cuts you off, or never stops listening	Raise or lower SILENCE_RMS in .env (see the mic test levels).
No voice, or "Voice error 403"	Her voice needs internet. Update it: pip install -U edge-tts
Whisper GPU errors at start	NVIDIA: re-run the nvidia-cublas / nvidia-cudnn pip line from Step 5. AMD: set WHISPER_DEVICE=cpu. She falls back to the CPU meanwhile.
AMD: very slow, ollama ps shows CPU	Your Radeon isn't supported by Ollama, or the driver is old: update Adrenalin.
Gmail opens a Google marketing page	Her Chrome window isn't signed in: sign in there once (Step 10).
"Web search isn't set up"	Add LANGSEARCH_API_KEY to .env (Step 7), then restart her.
She doesn't understand Tamil	Add GEMINI_API_KEY to .env. Free keys have a daily limit; after it, she uses English hearing.
Phone can't connect	Same Wi-Fi as the PC; allow Python through Windows Firewall; run python -m remote.pair again.

FAQ

Is it really free?

Yes. The model, Ollama, Whisper and the code are free. Her voice and the two keys use free services. You need your own PC, though.

Does it need ChatGPT or a subscription?

No. Her brain runs on your graphics card.

Can I change her name or voice?

Yes: ASSISTANT_NAME and JARVIS_VOICE in .env. Restart her after changes.

Can she talk back in Tamil?

She understands Tamil and Tanglish, but answers in English, on purpose: spoken Tamil from small local models wasn't good enough.

Does it work on AMD graphics?

Experimentally: her brain runs on a supported Radeon through Ollama, her hearing on the CPU, so each command is about 1-1.5 s slower. Not tested on AMD yet.

Will it work on my laptop?

Yes, with an NVIDIA RTX card with 8 GB VRAM (or a supported AMD Radeon) and 16 GB RAM. Keep it plugged in.

Command reference

Command	Purpose
python main.py	Start Keena (voice mode)
python main.py --text	Type instead of speaking
python main.py --mic-test	Show microphone levels
python main.py --brain-test	Check model speed and tool calling
python hud.py --demo (--big)	Preview the HUD (full screen with --big)
python -m remote.pair (--new)	Phone pairing QR code (--new makes a new key)
python -m brain.memory	List what she remembers
python -m brain.log	Show the conversation log
run_jarvis.cmd	Start Keena with automatic restart after a crash
scripts\autostart.ps1 (remove)	Start with Windows (remove to turn off)

Built with Qwen 3 8B (Ollama), faster-whisper, edge-tts, Playwright and PySide6. Released under the MIT License.